

HTTP Security

HTTP Threat Model

Eavesdropper


Listening on conversation (confidentiality)

Man-in-the-middle

Modifying content (integrity)

Impersonation

Bogus website (authentication, confidentiality)



HTTPS: Securing HTTP

HTTP sits on top of secure channel (SSL/TLS)


https:// vs. **http://**

TCP port 443 vs. 80

All (HTTP) bytes encrypted and authenticated

No change to HTTP itself!

Where to get the key???



Public Key Infrastructure

Public key certificate

Binding between **identity** and a **public key**

“Identity” is, for example, a domain name example.com

Digital signature to ensure integrity

Certificate authority

Issues public key **certificates** and verifies identities

Trusted parties (e.g., GoDaddy)


Preconfigured certificates in Web browsers

How to enable HTTPS for your server?

How to enable HTTPS for your server?

- Your Web Hosting Provider may offer HTTPS security or
- You can request a **SSL/TLS certificate** from Certificate Authorities and install it yourself.
- SSL/TLS certificates may need to be renewed periodically.

HTTP




Browser


User Id:
john.doe@emailaddress.com

→

Password:
@Apple123




Website's Server



With HTTP, hacker sees:

User Id: john.doe@emailaddress.com
Password: @Apple123


HTTPS




Browser

User Id:
john.doe@emailaddress.com

Password:
@Apple123



Website's Server




With HTTPS, hacker sees:

User Id: abErgdy#uwitWLqxytllqp


Password: xrtvxhj

HTTP vs. HTTPS



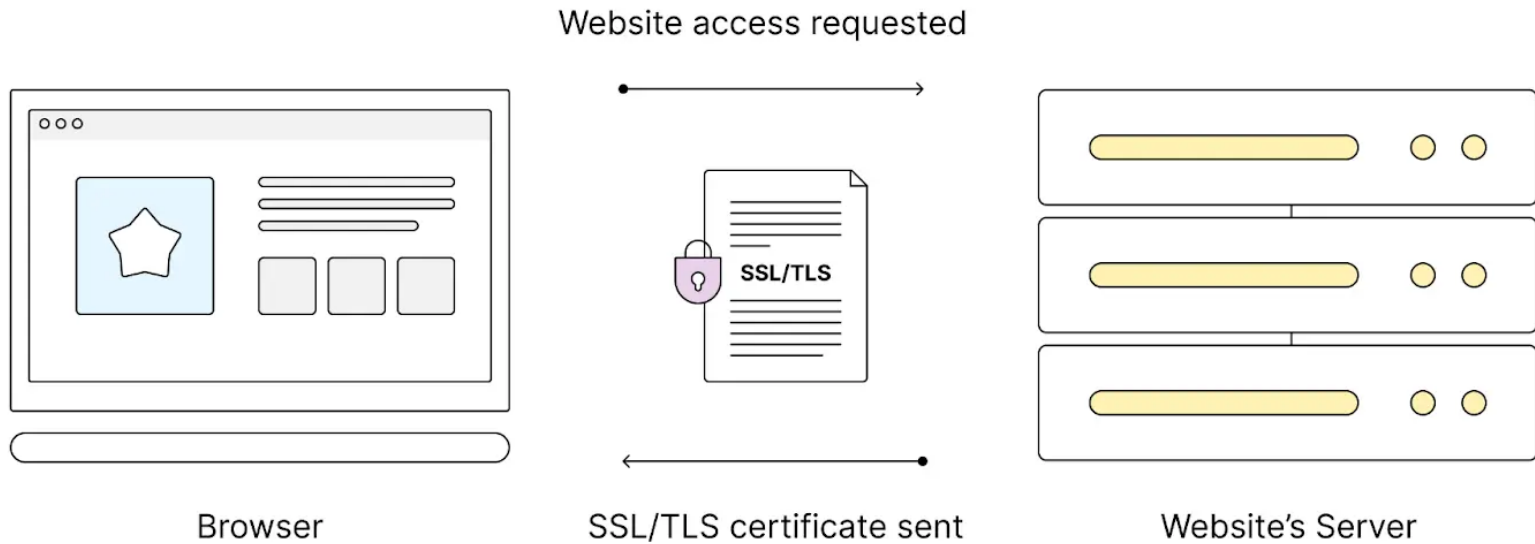
Browser

Website's Server




Browser

Website's Server



Browser ensures certificate is valid, not expired and matches the domain name.



Browser




Is the certificate valid?




Is it issued by a trusted
certificate authority?


Browser generates a “pre-master secret” (a temporary encryption key) and encrypts it with the server’s public key



Browser




The server uses its own private key to decrypt the pre-master secret




Website's Server


Using the pre-master secret, a **session key** is created that both the browser and server use to encrypt / decrypt messages (symmetric encryption).



Browser




Data is encrypted



Website's Server


HTTPS




Browser

User Id:
john.doe@emailaddress.com

Password:
@Apple123



Website's Server



With HTTPS, hacker sees:


User Id: abErgdy#uwitWLqxytllqp


Password: xrtvxhj

Large-Scale Web Applications


CPEN320

Web Application Architecture





Discuss solutions



Solutions?

1. Move to a stronger server
2. Add more servers

Scale-Up



Web Browser




Bigger Server



Bigger Storage System



Scale-Out



Scale-out architecture

- Expand capacity by adding more instances
- Contrast: **Scale-up architecture** - Switch to a bigger instance
 - Quickly hit limits on how big of single instances you can build
- Benefits of scale-out
 - Can scale to fit needs: Just add or remove instances
 - Natural redundancy make tolerating failures easier: One instance dies others keep working
- Challenge: Need to manage multiple instances and distribute work to them

Scale out web servers: Load Balancing

- Browsers want to speak HTTP to a web server - TCP/IP connect
- Use **load balancing** to distribute incoming HTTP requests across many front-end web servers
- HTTP redirection
 - Front-end machine accepts initial connections
 - Redirects them among an array of back-end machines
- DNS (Domain Name System) load balancing:
 - Specify multiple targets for a given domain name
 - Handles geographically distributed system
 - DNS servers rotate among those targets

Load-balancing switch ("Layer 4-7 Switch")

- Special load balancer network switch
 - Incoming packets pass through load balancer switch between Internet and web servers
 - Load balancer directs TCP connection request to one of the many web servers
 - Load balancer will send all packets for that connection to the same server.
- In some cases the switches are smart enough to inspect session cookies, so that the same session always goes to the same server.
- Stateless servers make load balancing easier (different requests from the same user can be handled by different servers).
- Can select web server based on random or on load estimates

nginx ("Engine X")

- Super efficient web server (i.e. speaks HTTP)
 - Handles 10s of thousands of HTTP connections
- Uses:
 - Load balancing - Forward requests to collection of front-end web servers
 - Handles front-end web servers coming and going (dynamic pools of server)
 - Fault tolerant - web server dies the load balance just quits using it
 - Handles some simple request - static files, etc.
 - DOS mitigation - request rate limits
- Popular approach to shielding Node.js web servers

Scale-out assumption: any web server will do

- Stateless servers make load balancing easier
 - Different requests from the same user can be handled by different servers
 - Requires database to be shared across web servers
- What about session state?
 - Accessed on every request so needs to be fast (memcache?)
- WebSockets bind browsers and web server
 - Can not load balance each request


Scale-out storage system

- Traditionally Web applications have started off using relational databases
- A single database instance doesn't scale very far.
- **Data sharding** - Spread database over scale-out instances
 - Each piece is called **data shard**
 - Can tolerate failures by **replication** - place more than one copy of data (3 is common)
- Applications must partition data among multiple independent databases, which adds complexity.
 - Facebook initial model: One database instance per university
 - In 2009: Facebook had 4000 MySQL servers - Use hash function to select data shard

Memcache: main-memory caching system

- Key-value data stored in memory
- Used to cache results of recent database queries: hit and miss
- Much faster than databases:
 - 500-microsecond access time, vs. 10's of milliseconds
- Example: Facebook has over 200,000 memcache servers
 - Writes must still go to the DBMS, so no performance improvement for them
 - Cache misses still hurt performance
 - Must manage consistency in software (e.g., flush relevant memcache data when database gets modified)

Scale-out web architecture



Building this architecture is hard

- Large capital and time cost in buying and installing equipment
- Must become expert in datacenter management
- Figuring out the right number of different components hard
 - Depends on load demand

Scaling issues are hard for early web app

- Startup: Initially, can't afford expensive systems for managing large scale.
- But, application can suddenly become very popular ("**flash crowd**"); can be disastrous if application can not scale quickly.
- Many of the early web apps either lived or died by the ability to scale
 - Friendster vs. Facebook

Virtualization - Virtual and Physical machines


Virtual Machine Images (Disk Images)

Load Balancer

Web Server


Database Server

Memcache



Load balancer	1
Web Server	100
Database	50
Memcache	20

Physical Machines



Cloud Computing

- Idea: Use servers housed and managed by someone else
 - Use Internet to access them
- Virtualization is a key enabler

Specify your compute, storage, communication needs:
Cloud provider does the rest

- Examples:
 - Amazon EC2
 - Microsoft Azure
 - Google Cloud
 - Many others

Load balancer	1
Web Server	100
Database	50
Memcache	20

Cloud Computing Pros and Cons?

Cloud Computing Pros and Cons

- Key: Pay for the resources you use
 - No upfront capital cost
 - Need 1000s machines right now? Possible
 - Perfect fit for startups:
 - 1998 software startup: First purchase: server machines
 - 2024 software startup: No server machines
- Typically billing is on resources:
 - CPU core time, memory bytes, storage bytes, network bytes
- Runs extremely efficiently
 - Buy equipment in large quantities, get volume discounts
 - Hire a few experts to manage large numbers of machines
 - Place servers where space, electricity, and labor is cheap

Higher level interfaces to web app cloud services

- Managing a web app backend at the level of virtual machines requires system building skills
- If you don't need the full generality of virtual machines you can use some already scalable platform.
 - Don't need to manage OSes: **Container** systems like Docker/Kubernetes
 - Specify programs and dependencies that run as a process
 - Don't need to manage storage - Cloud database storage
 - Let the cloud run the database
 - Don't need to manage instances/load balancing: Serverless
 - Let the cloud run the scale-out compute infrastructure

Cloud Database Storage

- Rather than running database instances - Use cloud run databases
 - Cloud provider has experts at running large scale systems
- Example: Google Spanner, Amazon DynamoDB
 - You: define schema, provide data, access using queries
 - Cloud provider: runs storage services
- Features:
 - High Availability
 - High Performance
 - Global replication and region containment
 - Consistency
 - Security
 - Usage based pricing

Serverless Computing

What is serverless?

Serverless Computing

- Serverless computing is a cloud computing execution model where the cloud provider **dynamically manages the allocation and provisioning of servers.**
- The name "serverless" comes from the fact that the tasks of server management and capacity planning decisions **are hidden** from the developer or operator.
- **This doesn't mean there are no servers involved**; instead, it means that developers no longer need to be concerned about servers, as the infrastructure management is handled by the cloud provider.

Serverless architecture - Cloud provider

- Hand over web-servers to cloud infrastructure
- Developer just specifies code to run on each URL & HTTP verb
 - Like Node/Express handlers
- Examples:
 - Amazon Lambda Functions
 - Microsoft Azure Functions
 - Google Cloud Functions
- Cloud provides services only (no servers)
 - Handles all scale-out, reliability, infrastructure security, monitoring, etc.
 - Pay by the request - Enable to pack function execution into available server resources
- Web App backend: Schema specification for cloud storage, handler functions

Serverless approach: Amazon Lambda

- You provide pieces of code, URLs associated with each piece of code
- Amazon Lambda does the rest:
 - Allocate machines to run your code
 - Arrange for name mappings so that HTTP requests find their way to your code
 - Scale machine allocations up and down automatically as load changes
 - Lambda environment also includes a scalable storage system
- More constrained environment
 - Must use their infrastructure and supported environments: Python, JavaScript, Java, Go, ...

Content Distribution Network (CDN)

- Consider a read-only part of our web app (e.g. image, React JavaScript, etc.)
 - Browser needs to fetch but doesn't care where it comes from
- Content distribution network
 - Has many servers positions all over the world
 - You give them some content (e.g. image) and they give you an URL
 - You put that URL in your app (e.g. `<img src="..."`)
 - When user's browsers access that URL they are sent to the closest server (DNS trick)
- Benefits:
 - Faster serving of app contents
 - Reduce load on web app backend
- Only works on content that doesn't need to change often

Cloud Computing and Web Apps

- The pay-for-resources-used model works well for many web app companies
 - At some point if you use many resources it makes sense to build your own data center
- Many useful infrastructure services available:
 - Auto scaling (spinning up and down instances on load changes)
 - Geographic distribution (can have parts of the backend in different parts of the world)
 - Monitoring and reporting (what parts of web app is being used, etc.)
 - Fault handling (monitoring and mapping out failed servers)
- Cloud Application Programming Interfaces (APIs):
 - Analytics
 - Machine learning - Prediction, recommendation, etc.
 - Translation, image recognition, maps, etc.

Discussion

When does it make sense to build your own data center instead of using the cloud?

Discussion

1. **Consistent high demand** for computing resources
2. Strict data **security and privacy** requirements
3. Performance and latency concerns
4. Need for **customization** and **specialized** hardware
5. Predictable **long-term growth and resource needs**
6. High data transfer **costs** or large data volumes
7. **Control over compliance** and disaster recovery
8. Financial viability and **long-term commitment**